

# **Big Data Analytics**

## **Session 11**

### **Predicting Algae Bloom**

# Problem Description

- Harmful algae in rivers
  - A serious ecological problem
    - Strong impact on
      - river life forms and
      - water quality
- Objectives:
  - Monitor and perform an early forecast of algae blooms
    - to improve the quality of rivers
    - chemical monitoring is cheaper and easily automated than biological analysis (microscopic examination)
  - Provide a better understanding of the factors influencing the algae frequencies



# Data Collection



- Several water samples were collected in different European rivers at different times during a period of approximately 1 year.
- For each water sample,
  - different chemical properties were measured, as well as
  - the frequency of several harmful algae
- Some related characteristics were stored
  - the season of the year
  - the river size
  - the river speed
- Data was collected in the context of the ERUDIT research Network
  - available in the UCI machine learning repository
  - <http://archive.ics.uci.edu/ml/datasets/Coil+1999+Competition+Data>

# Data Description



- Two main datasets:
  - Training dataset
    - 200 observations
    - 11 predictors
      - Nominal (3): season, size, speed
      - Numerical (8): different chemical parameters measured in the water samples
        - » Maximum pH value, Minimum value of O<sub>2</sub> (Oxygen)
        - » Mean value of Cl, NO<sub>3</sub><sup>-</sup>, NH<sub>4</sub><sup>+</sup>, PO<sub>4</sub><sup>3-</sup>, PO<sub>4</sub>, chlorophyll
    - 7 responses
      - Seven frequency numbers of different harmful algae found in respective sample
  - Test dataset
    - 140 observations
    - 11 predictors
    - no responses

Goal: to predict the frequency of the seven algae for these 140 water samples

# Load the Data into R



- Download the data (in .txt form) to your working directory (getwd()) from <http://www.dcc.fc.up.pt/~ltorgo/DataMiningWithR/datasets2.html> or <https://archive.ics.uci.edu/ml/machine-learning-databases/coil-mld/coil.html>

– **Analysis.txt**: training data; **Eval.txt**: test data

<http://www.dcc.fc.up.pt/~ltorgo/DataMiningWithR/DataSets/Analysis.txt>

```
algae <- read.table('Analysis.txt', header=F, dec='.',  
  col.names=c('season','size','speed','mxPH','mnO2','Cl','NO3','NH4','oPO4',  
    'PO4','Chla','a1','a2','a3','a4','a5','a6','a7'),  
  na.strings=c('XXXXXXX'))
```

#header=F: indicates that the file to be read does not include a first line with variable names

#dec='.': the numbers use '.' to separate decimal places (e.g., 34.2)

#na.strings: unknown values are represented by XXXXXXXX

or library(DMwR)

```
or algae <- read.table('/Users/than/./Analysis.txt', header=F, dec='.', ... )
```

**/Users/than/./** is the directory where the training data is stored.

```
> head(algae)
```

	season	size	speed	mxPH	mnO2	Cl	NO3	NH4	oPO4	PO4	Chla	a1	a2	a3	a4	a5	a6	a7
1	winter	small	medium	8.00	9.8	60.800	6.238	578.000	105.000	170.000	50.0	0.0	0.0	0.0	0.0	34.2	8.3	0.0
2	spring	small	medium	8.35	8.0	57.750	1.288	370.000	428.750	558.750	1.3	1.4	7.6	4.8	1.9	6.7	0.0	2.1
3	autumn	small	medium	8.10	11.4	40.020	5.330	346.667	125.667	187.057	15.6	3.3	53.6	1.9	0.0	0.0	0.0	9.7
4	spring	small	medium	8.07	4.8	77.364	2.302	98.182	61.182	138.700	1.4	3.1	41.0	18.9	0.0	1.4	0.0	1.4
5	autumn	small	medium	8.06	9.0	55.350	10.416	233.700	58.222	97.580	10.5	9.2	2.9	7.5	0.0	7.5	4.1	1.0
6	winter	small	high	8.25	13.1	65.750	9.248	430.000	18.250	56.667	28.4	15.1	14.6	1.4	0.0	22.5	12.6	3.9

# **Descriptive Data Analysis**


Data Visualisation and Summarisation

# Data Visualisation and Summarisation



- Use `summary(algae)`
  - Notice the difference that nominal and numerical variables are presented
    - Nominal: frequency counts
    - Numerical: 5 number summary

Histogram of `algae$mxPH`




- Use graphs to check the shape of distribution
- ```
> hist(algae$mxPH) # the shape suggests that mxPH is nearly normal distributed
```

# Data Visualisation and Summarisation


- Or boxplot

```
> boxplot(algae$oPO4, ylab='Orthophosphate (oPO4)')  
> abline(h=mean(algae$oPO4, na.rm=T), lty=2)
```






# Boxplots and Outliers



In R, all the suspected outliers and outliers are unfilled circles.

# Boxplots and Outliers



data is 1,2,3,4,5,6,7,8,x

$Q1=3$ ,  $Q2=5$ ,  $Q3=7$


$IQR = 7 - 3 = 4$

$1.5 * IQR = 6$


$3 * IQR = 12$

inner fence =  $Q3 + 1.5 * IQR = 13$


outer fence =  $Q3 + 3 * IQR = 19$




values1 = c(1:8, 10)



values2 = c(1:8, 13)



values3 = c(1:8, 14)




values4 = c(1:8, 20)

# Outliers – A Remark




- Outliers are not necessarily "bad" data-points
- They may well be the most important, most information rich, part of the dataset
- Under no circumstances should they be automatically removed from the dataset
- Outliers may deserve special consideration
  - they may be the key to the phenomenon under study or the result of human blunders


# rug and jitter



```
> values <- c(1,1,1,1,1,1,2)
> boxplot(values)
```



```
> rug(values, side=2)
```



```
> values = c(1,1,1,1,1,1,2)
> boxplot(values, ylim=c(0.7,2.2))
> set.seed(7)
> rug(jitter(values), side=2)
```


```
> set.seed(7)
> jitter(values)
```

```
[1] 1.1955637 0.9590982 0.8462791 0.8278995 0.8974998 1.1168042 1.9360249
```

# Data Visualisation and Summarisation

- Or boxplot

```
> boxplot(algae$oPO4, ylab='Orthophosphate (oPO4)')  
> abline(h=mean(algae$oPO4, na.rm=T), lty=2)  
> rug(jitter(algae$oPO4), side=2) #side=2 - left, 3-up, 4-right, 1-bottom
```



**jitter:**

Add a small amount of noise to a numeric vector.

**rug:**

Adds a set of tick marks along the base of a plot.

# Data Visualisation and Summarisation



- Detect outliers with graphics

```
>plot(algae$NH4,xlab='')
>abline(h=mean(algae$NH4,na.rm=T),lty=1,col="red")
>abline(h=mean(algae$NH4,na.rm=T)+sd(algae$NH4,na.rm=T),lty=2,col="blue")
>abline(h=median(algae$NH4,na.rm=T),lty=3,col="green")
>identify(algae$NH4)
```

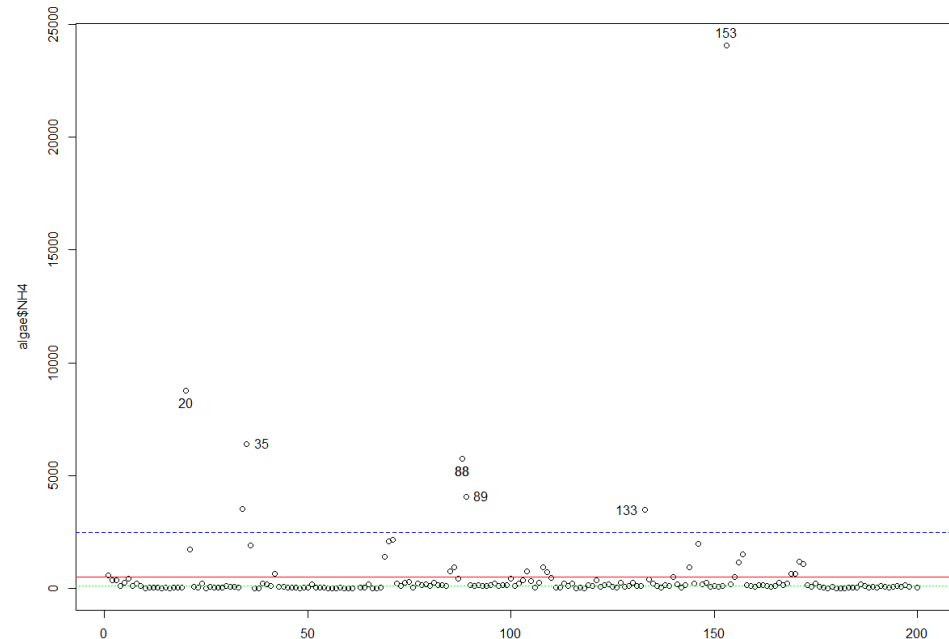
# identify is interactive: when a user click on the plotted dots with the left mouse, the row number of that observation will be shown. Click right mouse to finish interaction.

identify might not work in RStudio, though. Try the original R tool instead.

- Detect outliers without graphics

```
> algae[algae$NH4 >19000,]
```

|      | season | size   | speed | mxPH | mnO2 | Cl     | NO3   | NH4   | oPO4 | PO4 | Chla | a1  | a2 | a3 | a4  | a5  | a6   | a7   |
|------|--------|--------|-------|------|------|--------|-------|-------|------|-----|------|-----|----|----|-----|-----|------|------|
| NA   | <NA>   | <NA>   | <NA>  | NA   | NA   | NA     | NA    | NA    | NA   | NA  | NA   | NA  | NA | NA | NA  | NA  | NA   | NA   |
| 153  | autumn | medium | high  | 7.3  | 11.8 | 44.205 | 45.65 | 24064 | 44   | 34  | 53.1 | 2.2 | 0  | 0  | 1.2 | 5.9 | 77.6 | 0    |
| NA.1 | <NA>   | <NA>   | <NA>  | NA   | NA   | NA     | NA    | NA    | NA   | NA  | NA   | NA  | NA | NA | NA  | NA  | NA   | 14NA |



# **Data Preprocessing**

## Dealing with Missing Values

# Dealing With Unknown Values



- Unknown (missing) values
  - are common in real-world problems
  - may preclude the use of certain statistical learning approaches
    - E.g., `randomForest()`, `mean()`
- Solutions
  - Remove the cases with unknowns
  - Fill in the unknown values by exploring the most frequent value
  - Fill in the unknown values by exploring the correlations between variables
  - Fill in the unknown values by exploring the similarity between cases
  - Use tools that are able to handle these values



# Removing the Obs. with Unknown Values

- Before removing them, check/count them first

```
> library(DMwR)
> data(algae)           # load fresh data again before we try different ways of dealing with unknown values
> algae[!complete.cases(algae),]    # check whether each obs is complete or not
```

|     | season | size   | speed  | mxPH | mnO2 | Cl    | NO3   | NH4 | oPO4    | PO4     | Chla  | a1   | a2   | a3  | a4   | a5  | a6  | a7  |
|-----|--------|--------|--------|------|------|-------|-------|-----|---------|---------|-------|------|------|-----|------|-----|-----|-----|
| 28  | autumn | small  | high   | 6.80 | 11.1 | 9.000 | 0.630 | 20  | 4.000   | NA      | 2.70  | 30.3 | 1.9  | 0.0 | 0.0  | 2.1 | 1.4 | 2.1 |
| 38  | spring | small  | high   | 8.00 | NA   | 1.450 | 0.810 | 10  | 2.500   | 3.000   | 0.30  | 75.8 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 48  | winter | small  | low    | NA   | 12.6 | 9.000 | 0.230 | 10  | 5.000   | 6.000   | 1.10  | 35.5 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 55  | winter | small  | high   | 6.60 | 10.8 | NA    | 3.245 | 10  | 1.000   | 6.500   | NA    | 24.3 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 56  | spring | small  | medium | 5.60 | 11.8 | NA    | 2.220 | 5   | 1.000   | 1.000   | NA    | 82.7 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 57  | autumn | small  | medium | 5.70 | 10.8 | NA    | 2.550 | 10  | 1.000   | 4.000   | NA    | 16.8 | 4.6  | 3.9 | 11.5 | 0.0 | 0.0 | 0.0 |
| 58  | spring | small  | high   | 6.60 | 9.5  | NA    | 1.320 | 20  | 1.000   | 6.000   | NA    | 46.8 | 0.0  | 0.0 | 28.8 | 0.0 | 0.0 | 0.0 |
| 59  | summer | small  | high   | 6.60 | 10.8 | NA    | 2.640 | 10  | 2.000   | 11.000  | NA    | 46.9 | 0.0  | 0.0 | 13.4 | 0.0 | 0.0 | 0.0 |
| 60  | autumn | small  | medium | 6.60 | 11.3 | NA    | 4.170 | 10  | 1.000   | 6.000   | NA    | 47.1 | 0.0  | 0.0 | 0.0  | 0.0 | 1.2 | 0.0 |
| 61  | spring | small  | medium | 6.50 | 10.4 | NA    | 5.970 | 10  | 2.000   | 14.000  | NA    | 66.9 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 62  | summer | small  | medium | 6.40 | NA   | NA    | NA    | NA  | NA      | 14.000  | NA    | 19.4 | 0.0  | 0.0 | 2.0  | 0.0 | 3.9 | 1.7 |
| 63  | autumn | small  | high   | 7.83 | 11.7 | 4.083 | 1.328 | 18  | 3.333   | 6.667   | NA    | 14.4 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 116 | winter | medium | high   | 9.70 | 10.8 | 0.222 | 0.406 | 10  | 22.444  | 10.111  | NA    | 41.0 | 1.5  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 161 | spring | large  | low    | 9.00 | 5.8  | NA    | 0.900 | 142 | 102.000 | 186.000 | 68.05 | 1.7  | 20.6 | 1.5 | 2.2  | 0.0 | 0.0 | 0.0 |
| 184 | winter | large  | high   | 8.00 | 10.9 | 9.055 | 0.825 | 40  | 21.083  | 56.091  | NA    | 16.8 | 19.6 | 4.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 199 | winter | large  | medium | 8.00 | 7.6  | NA    | NA    | NA  | NA      | NA      | NA    | 0.0  | 12.5 | 3.7 | 1.0  | 0.0 | 0.0 | 4.9 |

```
>
> nrow(algae[!complete.cases(algae),])
[1] 16
>
> algae <- na.omit(algae)
[1] 184
```

# Removing the Obs. with Unknown Values

- Probably think twice before removing so many observations

```
> library(DMwR)
> data(algae) # load fresh data again before we try different ways of dealing with unknown values
> algae[!complete.cases(algae),] # check whether each observation is complete or not
```

|     | season | size   | speed  | mxPH | mnO2 | Cl    | NO3   | NH4 | oPO4    | PO4     | Chla  | a1   | a2   | a3  | a4   | a5  | a6  | a7  |
|-----|--------|--------|--------|------|------|-------|-------|-----|---------|---------|-------|------|------|-----|------|-----|-----|-----|
| 28  | autumn | small  | high   | 6.80 | 11.1 | 9.000 | 0.630 | 20  | 4.000   | NA      | 2.70  | 30.3 | 1.9  | 0.0 | 0.0  | 2.1 | 1.4 | 2.1 |
| 38  | spring | small  | high   | 8.00 | NA   | 1.450 | 0.810 | 10  | 2.500   | 3.000   | 0.30  | 75.8 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 48  | winter | small  | low    | NA   | 12.6 | 9.000 | 0.230 | 10  | 5.000   | 6.000   | 1.10  | 35.5 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 55  | winter | small  | high   | 6.60 | 10.8 | NA    | 3.245 | 10  | 1.000   | 6.500   | NA    | 24.3 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 56  | spring | small  | medium | 5.60 | 11.8 | NA    | 2.220 | 5   | 1.000   | 1.000   | NA    | 82.7 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 57  | autumn | small  | medium | 5.70 | 10.8 | NA    | 2.550 | 10  | 1.000   | 4.000   | NA    | 16.8 | 4.6  | 3.9 | 11.5 | 0.0 | 0.0 | 0.0 |
| 58  | spring | small  | high   | 6.60 | 9.5  | NA    | 1.320 | 20  | 1.000   | 6.000   | NA    | 46.8 | 0.0  | 0.0 | 28.8 | 0.0 | 0.0 | 0.0 |
| 59  | summer | small  | high   | 6.60 | 10.8 | NA    | 2.640 | 10  | 2.000   | 11.000  | NA    | 46.9 | 0.0  | 0.0 | 13.4 | 0.0 | 0.0 | 0.0 |
| 60  | autumn | small  | medium | 6.60 | 11.3 | NA    | 4.170 | 10  | 1.000   | 6.000   | NA    | 47.1 | 0.0  | 0.0 | 0.0  | 0.0 | 1.2 | 0.0 |
| 61  | spring | small  | medium | 6.50 | 10.4 | NA    | 5.970 | 10  | 2.000   | 14.000  | NA    | 66.9 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 62  | summer | small  | medium | 6.40 | NA   | NA    | NA    | NA  | NA      | 14.000  | NA    | 19.4 | 0.0  | 0.0 | 2.0  | 0.0 | 3.9 | 1.7 |
| 63  | autumn | small  | high   | 7.83 | 11.7 | 4.083 | 1.328 | 18  | 3.333   | 6.667   | NA    | 14.4 | 0.0  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 116 | winter | medium | high   | 9.70 | 10.8 | 0.222 | 0.406 | 10  | 22.444  | 10.111  | NA    | 41.0 | 1.5  | 0.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 161 | spring | large  | low    | 9.00 | 5.8  | NA    | 0.900 | 142 | 102.000 | 186.000 | 68.05 | 1.7  | 20.6 | 1.5 | 2.2  | 0.0 | 0.0 | 0.0 |
| 184 | winter | large  | high   | 8.00 | 10.9 | 9.055 | 0.825 | 40  | 21.083  | 56.091  | NA    | 16.8 | 19.6 | 4.0 | 0.0  | 0.0 | 0.0 | 0.0 |
| 199 | winter | large  | medium | 8.00 | 7.6  | NA    | NA    | NA  | NA      | NA      | NA    | 0.0  | 12.5 | 3.7 | 1.0  | 0.0 | 0.0 | 4.9 |

```
>
> manyNAs(algae) #returns the row numbers that have more than 20% of the columns with an NA. In this case, 18*20% = 3.6 columns.
[1] 62 199
> algae <- algae[-c(62,199),]
> algae <- algae[-manyNAs(algae),] # the last two commands have the same effect
```

# Filling with Most Frequent Values



- Several alternatives can be chosen, with different trade-offs between
  - the level of approximation, and
  - the computational complexity of the method
- First alternative (simplest and fastest)
  - Use some statistics of centrality to fill in the unknown values
    - mean, median, mode, etc
      - choose mean if the distribution is nearly normal
      - choose median if not
  - For example,

This method is simple, fast, thus appealing for large dataset. However, it may introduce a large bias in the data.

|    | season | size  | speed | mxPH | mnO2 | Cl    | NO3   | NH4 | oPO4  | PO4   | Chla | a1   | a2  | a3  | a4  | a5  | a6  | a7  |
|----|--------|-------|-------|------|------|-------|-------|-----|-------|-------|------|------|-----|-----|-----|-----|-----|-----|
| 48 | winter | small | low   | NA   | 12.6 | 9.000 | 0.230 | 10  | 5.000 | 6.000 | 1.10 | 35.5 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |

Recall that the mxPH is nearly normal distributed, we could use its mean value to fill in the hole.

```
>algae[48,'mxPH'] <- mean(algae$mxPH,na.rm=T)
```

```
#calculate the mean of the mxPH column while ignoring any NA values in this column
```

# Filling by Exploring Correlations

- An alternative to get less biased estimators for unknowns:
  - to explore the relationships between variables

```
> cor(algae[,4:18], use="complete.obs")           #disregard obs with NAs
> symnum(cor(algae[,4:18], use="complete.obs"))    #Symbolically encode a
given numeric or logical vector or array
```

|      | mP | mO | C1 | NO | NH | o | P | Ch | a1 | a2 | a3 | a4 | a5 | a6 | a7 |
|------|----|----|----|----|----|---|---|----|----|----|----|----|----|----|----|
| mxPH | 1  |    |    |    |    |   |   |    |    |    |    |    |    |    |    |
| mnO2 |    | 1  |    |    |    |   |   |    |    |    |    |    |    |    |    |
| C1   |    |    | 1  |    |    |   |   |    |    |    |    |    |    |    |    |
| NO3  |    |    |    | 1  |    |   |   |    |    |    |    |    |    |    |    |
| NH4  |    |    |    |    | 1  |   |   |    |    |    |    |    |    |    |    |
| oPO4 |    |    |    |    |    | 1 |   |    |    |    |    |    |    |    |    |
| PO4  |    |    |    |    |    | * | 1 |    |    |    |    |    |    |    |    |
| Chla |    |    |    |    |    |   |   | 1  |    |    |    |    |    |    |    |
| a1   |    |    |    |    |    |   |   |    | 1  |    |    |    |    |    |    |
| a2   |    |    |    |    |    |   |   |    |    | 1  |    |    |    |    |    |
| a3   |    |    |    |    |    |   |   |    |    |    | 1  |    |    |    |    |
| a4   |    |    |    |    |    |   |   |    |    |    |    | 1  |    |    |    |
| a5   |    |    |    |    |    |   |   |    |    |    |    |    | 1  |    |    |
| a6   |    |    |    |    |    |   |   |    |    |    |    |    |    | 1  |    |
| a7   |    |    |    |    |    |   |   |    |    |    |    |    |    |    | 1  |

NH4 and NO3 are positively correlated (0.72)

PO4 and oPO4 are highly correlated (above 0.9)

According to the domain expert, this was expected because the value of the total PO4 includes the value of oPO4

We will find the form of the linear correlation between these variables.

```
[1] 0 '' 0.3 '' 0.6 ' 0.8 '+' 0.9 '*' 0.95 'B' 1
```

# How to Find Linear Relationship



- Find linear relationship between  $\text{PO}_4$  and  $\text{oPO}_4$

```
> data(algae)
> algae <- algae[-manyNAs(algae),]
> lm(PO4 ~ oPO4, data=algae)
```

Call:

```
lm(formula = PO4 ~ oPO4, data = algae)
```

Coefficients:

| (Intercept) | oPO4  |
|-------------|-------|
| 42.897      | 1.293 |

The linear model we have obtained is  $\text{PO}_4 = 42.897 + 1.293 \times \text{oPO}_4$

- With this formula, we can fill in the unknown values of these unknowns, provided they are **not** both unknown.
  - Remove the observations with both unknown (sample 62, 199)
  - We have a single observation with an unknown value on  $\text{PO}_4$  (sample 28)

# Use the Linear Model to Predict



- Use  $PO_4 = 42.897 + 1.293 \times oPO_4$  to predict the unknown  $PO_4$  at sample 28

```
> algae[28, 'PO4'] <- 42.897 + 1.293 * algae[28, 'oPO4']
```

```
> algae[28,]
```

|    | season | size  | speed | mxPH | mnO2 | Cl | NO3  | NH4 | oPO4 | PO4    | Chla | a1   | a2  | a3 | a4 | a5  | a6  | a7  |
|----|--------|-------|-------|------|------|----|------|-----|------|--------|------|------|-----|----|----|-----|-----|-----|
| 28 | autumn | small | high  | 6.8  | 11.1 | 9  | 0.63 | 20  | 4    | 48.069 | 2.7  | 30.3 | 1.9 | 0  | 0  | 2.1 | 1.4 | 2.1 |

- This can be generalised to fill all missing  $PO_4$  values (if any)

```
> data(algae)
```

```
> algae <- algae[-manyNAs(algae),] # delete both unknowns
```

```
> fillPO4 <- function(oP) {
```

```
  if (is.na(oP))
```

```
    return(NA)          #if oPO4's value not available
```

```
  else
```

```
    return(42.897 + 1.293 * oP) #else return the result derived by linear model
```

```
}
```

```
> algae[is.na(algae$PO4), 'PO4'] <-
```

```
  sapply(algae[is.na(algae$PO4), 'oPO4'], fillPO4)
```


#This function is applied to all samples with unknown value on the variable  $PO_4$

# Filling by Exploring Correlations

- For other observations with unknown values, we can explore the correlations between the variables and the nominal variables of this problem.

- E.g., mxPH and season

```
> histogram(~ mxPH | season, data=algae)
```




The values of mxPH are not seriously influenced by the season of the year when the samples were collected.

# Filling by Exploring Correlations

- If we try the same using the size of the river

```
> histogram(~ mxPH | size, data=algae)
```



What tendency can you observe?



# Filling by Exploring Similar Cases



- Another alternative is to use the similarities between the rows to fill in the unknown values
    - If two water samples are similar, and one of them has an unknown value
      - It's very probable that this value is similar to the value of the other sample
    - How to define distance? Which distance can you think of?
      - Euclidean distance!
    - Approach:
      - Find ten most similar cases of any water sample with some unknown value
      - Use their values to fill in the unknown
        - The median of the values of the ten nearest neighbours
          - > algae <- knnImputation(algae, k=10, meth='median')
        - The weighted average of the values of the neighbours
          - » The further a neighbour is, the less weight it has (usual weight:  $1/d$ )
- using this function will replace all missing values in the df

# Obtaining Prediction Models

Multiple Linear Regression

Regression Trees

# Multiple Linear Regression



- The implementation of linear regression in R is not able to use datasets with unknown values

- Use the knn-preprocessed technique to fill in the unknowns.

```
> data(algae)
> algae <- algae[-manyNAs(algae), ]
> clean.algae <- knnImputation(algae, k = 10)
```

- Multiple linear regression

```
> lm.a1 <- lm(a1 ~ ., data=clean.algae[,1:12]) # here consider a1 with other 11 predictors
.....
Coefficients:
```

|              | Estimate  | Std. Error | t value | Pr(> t )  |
|--------------|-----------|------------|---------|-----------|
| (Intercept)  | 42.942055 | 24.010879  | 1.788   | 0.07537 . |
| seasonspring | 3.726978  | 4.137741   | 0.901   | 0.36892   |
| seasonsummer | 0.747597  | 4.020711   | 0.186   | 0.85270   |
| seasonwinter | 3.692955  | 3.865391   | 0.955   | 0.34065   |
| sizemedium   | 3.263728  | 3.802051   | 0.858   | 0.39179   |
| sizesmall    | 9.682140  | 4.179971   | 2.316   | 0.02166 * |
| speedlow     | 3.922084  | 4.706315   | 0.833   | 0.40573   |
| speedmedium  | 0.246764  | 3.241874   | 0.076   | 0.93941   |

Nominal variables are encoded by dummy variables

Erh, where are **seasonautumn**, **sizelarge** and **speedhigh**?

these are the base case (n-1 variables)

# Measures Explained

Coefficients:

|             | Estimate         | Std. Error      | t value       | Pr(> t ) ← p value |
|-------------|------------------|-----------------|---------------|--------------------|
| (Intercept) | 42.942055        | 24.010879       | 1.788         | 0.07537 .          |
| .....       |                  |                 |               |                    |
| mxPH        | -3.589118        | 2.703528        | -1.328        | 0.18598            |
| mnO2        | 1.052636         | 0.705018        | 1.493         | 0.13715            |
| Cl          | -0.040172        | 0.033661        | -1.193        | 0.23426            |
| <u>NO3</u>  | <u>-1.511235</u> | <u>0.551339</u> | <u>-2.741</u> | <u>0.00674 **</u>  |
| NH4         | 0.001634         | 0.001003        | 1.628         | 0.10516            |
| <u>oPO4</u> | <u>-0.005435</u> | <u>0.039884</u> | <u>-0.136</u> | <u>0.89177</u>     |
| PO4         | -0.052241        | 0.030755        | -1.699        | 0.09109 .          |
| Chla        | -0.088022        | 0.079998        | -1.100        | 0.27265            |

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

- t test: to see whether each coefficient is statistically significant (hypothesis  $H_0: \beta_i=0$ )
- $\text{Pr}(>|t|)$ : a value 0.0001 means that we are 99.99% confident that the coefficient is not null
  - Large value → insignificant factor, small value → significant factor (notice those with \*'s by R)

# Measures Explained



- $R^2$  coefficients (multiple and adjusted)
  - Degree of fit of the model
    - pve: proportion variance explained (the smaller, the lack of fit)
    - The adjusted coefficient is more demanding, as it takes into account the number of parameters in the model

Multiple R-squared: 0.3731, **Adjusted R-squared: 0.3215**

- F-statistics and p-value
  - To test  $H_0: \beta_1 = \beta_2 = \dots = \beta_m = 0$   
(the target variable does not depend on any of the predictors)
    - p-value: 0.0001 means that we are 99.99% confident that the null hypothesis is not true.
      - If p value is too high ( $>0.1$ ), it makes no sense to look at the t-test on individual coefficients

F-statistic: 7.223 on 15 and 182 DF, p-value: 2.444e-12

# Simply the Linear Model

- Some predictors have a small significance, we could eliminate them from the model

```
> anova(lm.a1) # ANOVA: analysis of variance
Analysis of Variance Table
```

```
Response: a1
```

|           | Df  | Sum Sq | Mean Sq | F value | Pr(>F)    |     |
|-----------|-----|--------|---------|---------|-----------|-----|
| season    | 3   | 85     | 28.2    | 0.0905  | 0.9651944 |     |
| size      | 2   | 11401  | 5700.7  | 18.3088 | 5.69e-08  | *** |
| speed     | 2   | 3934   | 1967.2  | 6.3179  | 0.0022244 | **  |
| mxPH      | 1   | 1329   | 1328.8  | 4.2677  | 0.0402613 | *   |
| mnO2      | 1   | 2287   | 2286.8  | 7.3444  | 0.0073705 | **  |
| Cl        | 1   | 4304   | 4304.3  | 13.8239 | 0.0002671 | *** |
| NO3       | 1   | 3418   | 3418.5  | 10.9789 | 0.0011118 | **  |
| NH4       | 1   | 404    | 403.6   | 1.2963  | 0.2563847 |     |
| oPO4      | 1   | 4788   | 4788.0  | 15.3774 | 0.0001246 | *** |
| PO4       | 1   | 1406   | 1405.6  | 4.5142  | 0.0349635 | *   |
| Chla      | 1   | 377    | 377.0   | 1.2107  | 0.2726544 |     |
| Residuals | 182 | 56668  | 311.4   |         |           |     |

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

It will give us the **reduction** in the **residual sum of squares** when adding each variable in turn

season contributes the least to the reduction of the fitting error of the model

# Update the Model



- Remove season from the model

```
> lm2.a1 <- update(lm.a1, . ~ . - season)
> summary(lm2.a1)
```

keep all x, minus season

Call:

```
lm(formula = a1 ~ size + speed + mxPH + mnO2 + Cl + NO3 + NH4 +
    oPO4 + PO4 + Chla, data = clean.algae[, 1:12])
```

Coefficients:

|             | Estimate   | Std. Error | t value | Pr(> t ) |
|-------------|------------|------------|---------|----------|
| (Intercept) | 44.9532874 | 23.2378377 | 1.934   | 0.05458  |

.....

|             |            |           |        |         |
|-------------|------------|-----------|--------|---------|
| speedmedium | -0.2976867 | 3.1818585 | -0.094 | 0.92556 |
| mxPH        | -3.2684281 | 2.6576592 | -1.230 | 0.22033 |
| mnO2        | 0.8011759  | 0.6589644 | 1.216  | 0.22561 |
| Cl          | -0.0381881 | 0.0333791 | -1.144 | 0.25407 |

.....

Multiple R-squared: 0.3682, **Adjusted R-squared: 0.3272**

F-statistic: 8.984 on 12 and 185 DF, p-value: 1.762e-13

This fit has improved a bit,  
but still not too impressive

Previously, **adjusted**  
**R-squared: 0.3215**

# Further AVONA Analysis

- Comparison between the two models

```
> anova(lm.a1, lm2.a1)
```

Analysis of Variance Table

Model 1:  $a1 \sim \text{season} + \text{size} + \text{speed} + \text{mxPH} + \text{mnO2} + \text{Cl} + \text{NO3} + \text{NH4} + \text{oPO4} + \text{PO4} + \text{Chla}$

Model 2:  $a1 \sim \text{size} + \text{speed} + \text{mxPH} + \text{mnO2} + \text{Cl} + \text{NO3} + \text{NH4} + \text{oPO4} + \text{PO4} + \text{Chla}$

|   | Res.Df | RSS   | Df | Sum of Sq | F      | Pr(>F) |
|---|--------|-------|----|-----------|--------|--------|
| 1 | 182    | 56668 |    |           |        |        |
| 2 | 185    | 57116 | -3 | -447.62   | 0.4792 | 0.6971 |

so much smaller than the first model

The second model is better, as it has a smaller sum of squares

However, with  $\text{Pr}(>F)=0.6971$ , it means that only with around 30% confidence we can say the two models are different

→ In other words, the difference between the two models are not significant

→ The second model is simpler

because there is one less feature (the season column)



# Automatic Model Simplification



- The `step` function will show you how to simplify the linear model step by step

```
> final.lm=step(lm.a1)
```

```
Start: AIC=1152.03
```

#AIC stands for Akaike Information Criterion

```
a1 ~ season + size + speed + mxPH + mnO2 + Cl + NO3 + NH4 + oPO4 + PO4 + Chla
```

```
#omit several steps in the middle, the last step is:
```

```
Step: AIC=1140.38
```

# step function use AIC to perform model search

```
a1 ~ size + mxPH + Cl + NO3 + PO4
```

|        | Df | Sum of Sq | RSS   | AIC    |
|--------|----|-----------|-------|--------|
| <none> |    |           | 58517 | 1140.4 |
| - mxPH | 1  | 784.1     | 59301 | 1141.0 |
| - Cl   | 1  | 835.6     | 59353 | 1141.2 |
| - NO3  | 1  | 1987.9    | 60505 | 1145.0 |
| - size | 2  | 2664.3    | 61181 | 1145.2 |
| - PO4  | 1  | 8575.8    | 67093 | 1165.5 |

AIC offers an estimate of the **relative information lost** when a given model is used to represent the process that generated.

It tells nothing about the absolute quality of a model, only the quality relative to other models. Thus, if all the candidate models fit poorly, AIC will not give any warning of that.

# Analyse the Final Model



```
> summary(final.lm)
```

Call:

```
lm(formula = a1 ~ size + mxPH + Cl + NO3 + PO4, data = clean.algae[, 1:12])
```

Residuals:

| Min     | 1Q      | Median | 3Q    | Max    |
|---------|---------|--------|-------|--------|
| -28.874 | -12.732 | -3.741 | 8.424 | 62.926 |

Coefficients:

|             | Estimate | Std. Error | t value | Pr(> t )     |
|-------------|----------|------------|---------|--------------|
| (Intercept) | 57.28555 | 20.96132   | 2.733   | 0.00687 **   |
| sizemedium  | 2.80050  | 3.40190    | 0.823   | 0.41141      |
| sizesmall   | 10.40636 | 3.82243    | 2.722   | 0.00708 **   |
| mxPH        | -3.97076 | 2.48204    | -1.600  | 0.11130      |
| Cl          | -0.05227 | 0.03165    | -1.651  | 0.10028      |
| NO3         | -0.89529 | 0.35148    | -2.547  | 0.01165 *    |
| PO4         | -0.05911 | 0.01117    | -5.291  | 3.32e-07 *** |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 17.5 on 191 degrees of freedom

Multiple R-squared: 0.3527, **Adjusted R-squared: 0.3324**

F-statistic: 17.35 on 6 and 191 DF, p-value: 5.554e-16

In multiple regression, the proportion of variance explained (PVE) is equal to (adjusted)  $R^2$ .

The PVE is still not very interesting (0.3324).

A sign that linearity assumption of this model is inadequate for the domain.

# Obtaining Prediction Models


Multiple Linear Regression

Regression Trees

# Build a Regression Tree

- It can be done in the same way as building a classification tree


```
> library(tree)
> data(algae)
> algae<-algae[-manyNAs(algae),]
> rt.a1<-tree(a1~.,algae[,1:12])
> text(rt.a1,pretty=0)
```



# Build the Tree using Train Part

- Now we randomly sample a train set and build the regression tree based on the set


```
> nrow(algae)
[1] 198
> set.seed(2)
> train.a1 <- sample(1:nrow(algae), nrow(algae)/2)
> rt.a1.train <- tree(a1~., algae[,1:12], subset=train)
> plot(rt.a1.train)
> text(rt.a1.train, pretty=0)
```



# Use CV to Check Whether to Prune

- Cross validation is used to see whether the tree `rt.al.train` needs to be pruned

```
> cv.rt.train <- cv.tree(rt.al.train)
> plot(cv.rt.train$size, cv.rt.train$dev, type='b')
```




The best tree (the one with the minimum MSE) is of the size 5

# Prune the Tree

- Prune the tree to be of size 5:


```
> prune.rt.a1 <- prune.tree(rt.a1.train,best=5)
> plot(prune.rt.a1)
> text(prune.rt.a1,pretty=0)
```



# Performance Evaluation – Regression Tree

- We use the test part to evaluate the performance

```
> rt.a1.test <- algae[-train,"a1"]  
> yhat.rt.a1.prune <- predict(prune.rt.a1,newdata=algae[-train,1:12])  
> mean((yhat.rt.a1.prune-rt.a1.test)^2)  
[1] 297.0548  
> plot(yhat.rt.a1.prune,rt.a1.test)  
> abline(0,1)
```





# Using Bagging



- Since the bagging/randomForest method requires no missing values, we start from the dataset `clean.algae`

```
> set.seed(20)
> bag.train <- sample(1:nrow(clean.algae), 99)
> bag.a1.train <- randomForest(a1 ~ ., clean.algae[1:12],
                               subset=bag.train, mtry=11, importance=T)
> bag.a1.train
```

Call:

```
randomForest(formula = a1 ~ ., data = clean.algae[, 1:12], mtry = 11,
              importance = T, subset = bag.train)
```

Type of random forest: regression

Number of trees: 500

No. of variables tried at each split: 11

when compared to the adjusted R-squared for the linear model, which is around 30%

Mean of squared residuals: 271.161

% Var explained: 42.9



This PVE is  
larger than the  
linear model<sup>11</sup>


# Performance Evaluation - Bagging



- How well does this bagged model perform on the test set?

```
> yhat.bag <- predict(bag.a1.train, newdata=clean.algae[-bag.train, 1:12])  
> clean.algae.test <- clean.algae[-bag.train, "a1"]  
> mean((yhat.bag-clean.algae.test)^2)  
[1] 279.9227  
> plot(yhat.bag, clean.algae.test)  
> abline(0, 1)
```

← This MSE is smaller than the best tree



This looks better than in the regression tree →


You may play with the number of trees in the bagging at home (ntree=i)

# Which Predictors are Important?

```
> importance(bag.a1.train)
```

|        | %IncMSE     | IncNodePurity |
|--------|-------------|---------------|
| season | -0.78908700 | 1168.1408     |
| size   | 2.04011486  | 439.9201      |
| speed  | 0.95764449  | 410.7687      |
| mxPH   | 3.15721118  | 1878.4553     |
| mnO2   | 0.07241614  | 1402.3103     |
| Cl     | 11.17164253 | 8361.5771     |
| NO3    | 4.13772152  | 1904.3340     |
| NH4    | 4.34200936  | 2552.3389     |
| oPO4   | 12.77819198 | 9659.2113     |
| PO4    | 13.97331416 | 9281.1592     |
| Chla   | 13.27516985 | 7345.5980     |

```
> varImpPlot(bag.a1.train)
```



# Using Random Forest



- Choose a smaller `mtry` value, usually  $p/3$  when building a random forest for regression trees
  - `mtry = 11/3  $\approx$  3 or 4`

```
> set.seed(20)
> rf.a1.train.4 <- randomForest(a1~.,clean.algae[,1:12], subset=bag.train,
                                mtry=4, importance=T)
> yhat.rf <- predict(rf.a1.train.4,newdata=clean.algae[-bag.train,1:12])
> mean((yhat.rf-clean.algae.test)^2)
[1] 273.3071
```

`mtry = 3` is slightly better

You may find the best `mtry` at home

```
> set.seed(20)
> rf.a1.train.3 <- randomForest(a1~.,clean.algae[,1:12],subset=bag.train,
                                mtry=3, importance=T)
> yhat.rf <- predict(rf.a1.train.3,newdata=clean.algae[-bag.train,1:12])
> mean((yhat.rf-clean.algae.test)^2)
[1] 272.3034
```

The PVE of `rf.a1.train.3` is 49.6% (use `summary()`), still not very fit

Probably try nonlinear models (polynomials, etc), something for you to try at home, too

# Prediction for New Test Set

# Prediction for the Algae



- We are given 140 test samples, whose algae levels are unknown.
- We will choose the best models to obtain these predictions.
  - To obtain unbiased estimates of MSE for a set of models
    - By means of a cross-validation experimental process
    - For simplicity, we only predict `a1`
- For `a1`, we have already shown that the randomForest model `rf.a1.train.3` is the best model
  - Use `rf.a1.train.3` to make the prediction

# Unknowns in the Test Data



- There are unknowns in the test data
- We could use `knnImputation()` as in the training dataset
  - Use other test cases to fill in the unknowns → not ideal
  - Use training data to find the neighbours instead
    - use `knnImputation()`, but with an extra argument

```
> clean.test.algae <- knnImputation(test.algae, k=10, distData=algae[, 1:11])
```

# The `distData` argument allows you to supply an extra set of data (i.e., the training dataset) where the ten nearest neighbours are to be found for each case with unknowns in the `test.algae` dataset.

# Make the Prediction



- Finally,...

```
> preds <- rep(0,140)
> preds <- predict(rf.a1.train.3, newdata=clean.test.algae, mtry=3, importance=T)
> preds
```

| 1         | 2         | 3         | 4         | 5         | 6         | 7         | 8         | 9         |
|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| 7.266943  | 10.458083 | 13.387457 | 13.542400 | 27.145823 | 33.591357 | 35.073133 | 37.611920 | 38.065740 |
| 10        | 11        | 12        | 13        | 14        | 15        | 16        | 17        | 18        |
| 36.190503 | 10.706703 | 15.288940 | 40.884010 | 38.163287 | 37.630820 | 26.044157 | 10.487700 | 20.337720 |
| 19        | 20        | 21        | 22        | 23        | 24        | 25        | 26        | 27        |
| 40.361043 | 54.538080 | 6.965607  | 4.724927  | 4.981443  | 11.896803 | 6.452217  | 5.023043  | 24.228200 |
| 28        | 29        | 30        | 31        | 32        | 33        | 34        | 35        | 36        |
| 43.114077 | 27.373763 | 23.633090 | 26.444843 | 20.911110 | 32.294507 | 38.157100 | 55.714590 | 35.624243 |
| 37        | 38        | 39        | 40        | 41        | 42        | 43        | 44        | 45        |
| 35.052197 | 51.597240 | 33.467427 | 39.437900 | 37.612970 | 16.618960 | 10.317370 | 9.975300  | 10.411817 |
| 46        | 47        | 48        | 49        | 50        | 51        | 52        | 53        | 54        |
| 3.337300  | 10.015007 | 5.438577  | 17.838527 | 31.355300 | 11.017717 | 3.678907  | 5.509753  | 3.913507  |
| 55        | 56        | 57        | 58        | 59        | 60        | 61        | 62        | 63        |
| 4.779007  | 12.729870 | 13.189073 | 11.902373 | 17.185123 | 14.290100 | 6.853557  | 21.050563 | 16.727573 |
| 64        | 65        | 66        | 67        | 68        | 69        | 70        | 71        | 72        |
| 8.964617  | 33.982597 | 27.070277 | 18.403937 | 40.085983 | 43.577550 | 4.610323  | 6.584670  | 4.338993  |
| .....     |           |           |           |           |           |           |           |           |





WORK  
HARD  
AND  
WAIT FOR  
CHRISTMAS